

Object Detection Extension

<https://github.com/aiyshawry/Object-Detection-Extension-in-Rapid-Miner>

[Python](#)

[PyTorch](#)

[ONNX Runtime](#)

[RapidMiner](#)

[Object Detection](#)

OVERVIEW

A RapidMiner Studio extension that brings transfer learning-based object detection into a no-code environment, supporting fine-tuning and inference with Faster R-CNN, FCOS, SSD, SSDLite, and RetinaNet models.

IMPACT

Developed two RapidMiner operators for object detection — `FineTuneObjectDetectionModel` and `ObjectDetectionModelInference` — enabling users to train and deploy detection models within a visual workflow without writing code.

TECHNICAL WALKTHROUGH

Object Detection Extension

Problem Statement Different object detection use cases require different trade-offs between accuracy, speed, and computational cost. A single model cannot satisfy all scenarios.

Approach Developed a unified object detection extension using PyTorch, supporting multiple architectures: Faster R-CNN FCOS SSD SSD Lite RetinaNet YOLO (extended work) Training Dataset annotated using XML bounding boxes Transfer learning using COCO pre-trained weights Fully configurable training pipeline: Model selection Batch size Epochs Learning rate Momentum Weight decay Evaluation metrics: IoU (Intersection over Union) Precision Recall (per class) Inference Users can choose any trained or pre-trained model Output includes: Image with bounding boxes, labels, and confidence scores Tabular output with detected objects Results also generated in base64 format for easy deployment and integration **Key Observations** Faster R-CNN performed best for accuracy SSD and SSD Lite were faster but had lower recall RetinaNet handled class imbalance well Model choice depends on application requirements

KEY DETAILS

- `FineTuneObjectDetectionModel` accepts image and annotation directories (YOLO/Pascal VOC) and outputs model weights and class list CSV.

- `ObjectDetectionModelInference` runs predictions on new images and returns bounding boxes and class labels as a `DataFrame`.

- Supports five architectures: Faster R-CNN, FCOS, SSD, SSDLite, and RetinaNet for varying speed-accuracy trade-offs.

- Configurable for CPU-only and GPU (CUDA cu118) environments; installed as a JAR extension in RapidMiner Studio.

- Enables domain-specific object detection with minimal ML expertise using a drag-and-drop process designer.