

Sentiment Analysis Extension

<https://github.com/aiyshwary/Video-based-and-audio-based-sentiment-analysis>

Python

Hugging Face

NLP

RapidMiner

Audio Processing

OVERVIEW

A multimodal sentiment analysis project that independently extracts and analyses sentiment from both the video (visual) and audio streams of a recording, combining insights from two separate NLP and signal processing pipelines.

IMPACT

Built a dual-stream sentiment analysis pipeline that separates video and audio from a recording and independently classifies sentiment from each modality, enabling richer multimodal emotional understanding.

TECHNICAL WALKTHROUGH

Sentiment Analysis Extension

Problem Statement

Understanding customer opinions and emotional tone from text data is crucial for decision-making in domains like customer feedback, reviews, and social media.

Approach

Built a sentiment analysis pipeline using:

NLP preprocessing (tokenization, cleaning, normalization) Pre-trained language models (via Hugging Face where applicable)

Supports classification into:

Positive

Negative

Neutral sentiments

Methodology

Text data is cleaned and vectorized

Models are trained or fine-tuned on labeled sentiment datasets

Predictions include:

Sentiment label

Confidence score

Applications

Customer feedback analysis

Product review monitoring

Social media sentiment tracking

Deployment

Entire system deployed through AI-Hub using RapidMiner

Images and data stored in AWS S3

Inference handled via automated workflows

Supports:

Real-time predictions

Error handling

Scalable deployment

Conclusion

This project delivers a complete AI ecosystem for:

Image Classification

Object Detection

Sentiment Analysis

The key strength lies in its flexibility, modularity, and deployability, allowing users to:

Choose models based on requirements

Fine-tune on custom datasets

Deploy seamlessly in production environments

It bridges the gap between research-level models and real-world industrial applications.

KEY DETAILS

- Extracts audio track from video files and applies speech-to-text followed by NLP-based sentiment classification.
- Processes video frames independently to infer visual sentiment cues from facial expressions and scene context.
- Outputs positive, neutral, or negative sentiment labels with confidence scores per modality.
- Designed for use cases including customer feedback analysis, social media monitoring, and presentation review.
- Implemented using Hugging Face pretrained models and RapidMiner for pipeline orchestration.