

# Video RAG Retrieval Project

<https://github.com/aiyshawry/Video-RAG-Retrieval>

Python

FAISS

OpenCLIP

ViT-GPT2

Sentence Transformers

RAG

## OVERVIEW

A semantic video search system that extracts frames, generates scene captions using ViT-GPT2, and embeds them with SentenceTransformers and OpenCLIP into a dual FAISS index, enabling natural language queries to retrieve the most relevant video scenes with timestamps.

## IMPACT

Built a Visual RAG system that makes unstructured video content semantically searchable: a text query retrieves the most relevant scenes with timestamps by matching against dual text and image embedding indexes via cosine similarity.

## TECHNICAL WALKTHROUGH

"I built a Visual Retrieval-Augmented Generation (RAG) system that retrieves relevant video scenes based on a user's text prompt by combining vision models and vector similarity search."

### Problem statement

"Searching videos using text is hard because video content is unstructured.

Traditional keyword search doesn't understand what's actually happening inside a video. My project solves this by converting both video scenes and user queries into embeddings and retrieving the most relevant video segments."

### High-level architecture

"The system has three main components

#### Scene understanding using vision models

#### Vector storage for text and image embeddings

Similarity-based retrieval using VDMS"

Video → Scene Description + Frames → Embeddings → Vector DB  
→ Retrieval → Video Clips

#### Input & preprocessing

"The input to the system is one or more videos. Each video is processed to

#### extract

- Scene-level textual descriptions using open-source vision models like

#### Video-LLaMA

Key video frames along with metadata like frame number, FPS, and timestamp"\*

#### Vision understanding

"I use vision-language models to understand the video content. These models analyze video scenes and generate natural language descriptions that summarize what's happening in each scene."

“For example, a scene might be described as ‘a person walking into a store and picking up an item’.”

### **Embedding strategy (very important for interviewers)**

“To support efficient retrieval, I store two types of embeddings in separate vector stores.”

### **Text embeddings**

Model: SentenceTransformer – all-MiniLM-L6-V2 / L12-V2

**Used for:**

### **Scene descriptions**

### **User queries**

“This converts textual descriptions into dense vectors that capture semantic meaning.”

### **Image embeddings**

Model: OpenCLIP – ViT-g-14

**Used for:**

### **Video frames**

“This helps capture visual semantics like objects, actions, and context directly from frames.”

### **Vector storage & retrieval**

“Both text and image embeddings are stored in VDMS, which allows fast vector similarity search. I maintain separate vector collections for text and image embeddings but link them using video metadata.”

### **Stored metadata includes:**

**Video name**

**Scene ID**

**Frame number**

**Timestamp**

**Query flow**

“When a user enters a text prompt:”

**Prompt is converted into a text embedding**

**Vector similarity search is performed against:**

**Scene text embeddings**

(Optionally) image embeddings

**The system retrieves:**

**Most relevant scenes**

**Corresponding frames and timestamps**

**Final output:**

**Video name**

**Scene description**

Exact time range to jump to in the video

question)

“Text embeddings capture semantic meaning, while image embeddings capture visual features. Using both improves retrieval accuracy because some queries are better answered by visual similarity and others by textual context.”

#### **Example:**

“‘a red car’ → image embeddings

‘a tense conversation’ → text embeddings”

#### **Technical depth**

**Efficient vector search using cosine similarity**

**Decoupling embeddings for scalability**

**Model choice based on speed vs accuracy**

**Frame sampling using FPS to avoid redundancy**

Closing line (impact)

“This approach makes video content searchable, explainable, and scalable, and it can be extended to surveillance, media indexing, or enterprise video search.”

#### **Short version**

“I built a Visual RAG system that enables semantic search over videos. It uses vision models to generate scene descriptions and extract frames, converts them into text and image embeddings using SentenceTransformers and OpenCLIP, stores them in VDMS, and retrieves relevant video segments based on user queries.”

#### **KEY DETAILS**

---

- Extracts frames at a configurable FPS and generates natural-language scene captions using ViT-GPT2 image captioning.
- Computes text embeddings with SentenceTransformers (MiniLM) and image embeddings with OpenCLIP for dual-modality indexing.
- Stores embeddings and metadata in two parallel FAISS indexes supporting independent text or image-based retrieval.
- Query interface accepts a natural language string and returns top-k matching scenes with frame references and timestamps.
- Modular architecture (preprocessing, embeddings, storage, retrieval) designed for extension to production vector databases.