

Video RAG Retrieval Project

<https://github.com/aiyshwary/Video-RAG-Retrieval>

OVERVIEW

A semantic video search system that extracts frames, generates scene captions via ViT-GPT2, and embeds them using OpenCLIP. Queries are matched against a dual FAISS index using cosine similarity, returning the most relevant results and frame references.

IMPACT

Enables semantic search over unstructured video content by leveraging vision-language models and vector search, making scenes retrievable by natural language queries.

TECHNICAL WALKTHROUGH

Implements a simple Video Retrieval-Augmented Generation (RAG) system for semantic video search using video frame embeddings and performing nearest-neighbor lookups.

- i) Extract frames and scene descriptions from videos using ViT-GPT2 image captioning.
- ii) Compute text embeddings with SentenceTransformers and image embeddings with OpenCLIP.
- iii) Store embeddings in FAISS indexes with associated metadata (video name, scene ID, frame number).
- iv) Retrieve relevant scenes or frames for a text query using cosine similarity.

Clone the repository, create a virtual environment, and install dependencies from requirements.txt. GPU acceleration is supported for faster computation.

Process a video and index its scenes: `extract_scenes('path/to/video.mp4', fps=0.5); retriever = VideoR retriever.index_scenes(scenes, video_name='demo').`
Query: `results = retriever.query('a person walking', k=3, use_text=True).`

/Video RAG Retrieval Project ■■■■ README.md ■■■■ requirements.txt ■■■■ src/video_rag/ ■■■■ __in
embeddings.py ■■■■ storage.py ■■■■ retrieval.py ■■■■ utils.py

- i) Scene detection is rudimentary; can be replaced with advanced shot-detection or segmentation logic
- ii) Metadata stored with embeddings can be extended to track file paths, timestamps, etc.
- iii) For production, swap FAISS for a proper vector database such as VDMS, Milvus, or Pinecone.

KEY DETAILS

1. Extracts frames and scene descriptions from videos using open-source models (ViT-GPT2 for captions)
2. Computes text embeddings with SentenceTransformers and image embeddings with OpenCLIP.
3. Stores embeddings in FAISS indexes with associated metadata.
4. Retrieves relevant scenes or frames for a text query.
5. Lightweight Python prototype; can be extended to production vector DBs like VDMS, Milvus, or Pinecone.